

Outliers Treatment

outliers are redundant observations. bcz they spoil the performance of the model.
Defination : points which are extremly far from the normal observations.

Where do they come from :

1. Human error: data entry error
2. system error : measurment error
3. use of psudo data: dummy dataset, fake data
4. sampling error: creating or processing dataset using wrong functions
5. data merging : mixing of datasets
6. Natural error: most of the errors might belong to this category

Treating of outliers:

1. dropping
2. winitializing (replacing the outliers with statistical computations(mean,median,mode etc..))
3. Thresold split
4. Normalizing
5. Tranformation:
 1. log transform
 2. Noramlization (0-1)
 3. standardization (-1 to 1)
 4. box-cox transform
 5. sqrt transform
 6. cbrr transform
 7. reciprocal transform

Algorithms that are not affected by outliers:

Decision Tree,
Random Forest
Adaboost, Xgboost
Naive Baiyes

Algorithms that are affected by outliers:

Linear Regression
Logistic Regression
K-NN
SVM
K-means clustering

Detecting outliers

check if the data is noramlly distributed
shapiro test, normality test, kstest


```
[2]: import numpy as np
```

```
[3]: array = np.random.randint(10,100,size=20)  
array
```

```
[3]: array([57, 82, 90, 73, 83, 56, 74, 90, 96, 40, 94, 77, 20, 79, 68, 64, 56,  
          88, 68, 57])
```

```
[4]: import seaborn as sns  
sns.kdeplot(array,fill=True)
```

```
[4]: <Axes: ylabel='Density'>
```




```
[8]: array2 = array.copy()
      array2[10]=200
      array2[12]=250
      array2[5]=-200
      array2

[8]: array([ 57,  82,  90,  73,  83, -200,  74,  90,  96,  40, 200,
           77, 250,  79,  68,  64,  56,  88,  68,  57])

[9]: sns.kdeplot(array2,fill=True)

[9]: <Axes: ylabel='Density'>
```



Z-score Normalization

```
z_score = (x_value - x_mean)/std
```

Idea is calculating `z_score` & setting up a threshold such that all the datapoints above the threshold or below the threshold are outliers & can be eliminated.

```
[10]: mean = np.mean(array2)
      std = np.std(array2)
      print('mean is ',mean, 'std is ',std)
```

```
outliers_list=[]  
threshold = 1.5  
  
for x_val in array2:  
    z_score = (x_val-mean)/std  
    z_score = np.abs(z_score)  
    print(z_score)  
  
    if z_score >= threshold:  
        outliers_list.append(x_val)
```

outliers_list

```
mean is 74.6 std is 79.22209287818646  
0.2221602505132264  
0.09340828714760667  
0.19439021919907326  
0.020196386410293244  
0.10603102865403999  
3.4662048176665907  
0.007573644903859923  
0.19439021919907326  
0.2701266682376732  
0.4367468561225929  
1.5828917849067388  
0.030294579615440048  
2.214028860228405  
0.05554006262830669  
0.08331009394245986  
0.13380105996819316  
0.23478299201965974  
0.1691447361862066  
0.08331009394245986  
0.2221602505132264
```

[10]: [-200, 200, 250]

[16]: np.where(array2<=min(outliers_list),np.mean(array2),array2)

```
[16]: array([ 57. , 82. , 90. , 73. , 83. , 74.6, 74. , 90. , 96. ,
          40. , 200. , 77. , 250. , 79. , 68. , 64. , 56. , 88. ,
          68. , 57. ])
```

Using IQR

Inter Quantile Range

```
[17]: array
```

```
[17]: array([57, 82, 90, 73, 83, 56, 74, 90, 96, 40, 94, 77, 20, 79, 68, 64, 56,
          88, 68, 57])
```

```
[18]: q1 = np.quantile(array,0.25)
      q1
```

```
[18]: 57.0
```

```
[19]: q2 = np.quantile(array,0.5)
      q2
```

```
[19]: 73.5
```

```
[20]: q2 = np.quantile(array,0.75)
      q2
```

```
[20]: 84.25
```

```
[22]: import pandas as pd
      pd.Series(array).describe()
```

```
[22]: count    20.000000
      mean     70.600000
      std     19.209099
      min     20.000000
      25%     57.000000
```

```
50%      73.500000
75%      84.250000
max      96.000000
dtype: float64
```

```
IQR
Q1 >> 1st quantile >> 25th percentile of the data
Q2 >> 2nd quantile >> 50th percentile of the data / Median
Q3 >> 3rd quantile >> 75th percentile of the data

IQR = Q3-Q1

lower_point = Q1 - 1.5*IQR
upper_bound = Q3 + 1.5*IQR
```

```
[23]: df = pd.read_csv('titanic.csv')
df
```

```
[23]:
```

	PassengerId	Survived	Pclass	Name	Gender	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000	NaN	S
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000	B42	S
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C148	C

890 891 0 3 Dooley, Mr. Patrick male 32.0 0 0 370376 7.7500 NaN Q

891 rows × 12 columns

```
[24]: df.dropna(subset=['Age'],inplace=True)
```

```
[26]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 714 entries, 0 to 890
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   PassengerId  714 non-null    int64  
 1   Survived     714 non-null    int64  
 2   Pclass       714 non-null    int64  
 3   Name         714 non-null    object  
 4   Gender       714 non-null    object  
 5   Age          714 non-null    float64 
 6   SibSp        714 non-null    int64  
 7   Parch        714 non-null    int64  
 8   Ticket       714 non-null    object  
 9   Fare         714 non-null    float64 
10   Cabin        185 non-null    object  
11   Embarked     712 non-null    object  
dtypes: float64(2), int64(5), object(5)
memory usage: 72.5+ KB
```

```
[27]: df.Age.describe()
```

```
[27]: count    714.000000
      mean     29.699118
      std     14.526497
      min       0.420000
      25%     20.125000
      50%     28.000000
      75%     38.000000
      max     80.000000
```


Name: Age, dtype: float64

```
[42]: import warnings
      warnings.filterwarnings("ignore")
```

```
[43]: Q1 = df.Age.describe()[4]
      Q2 = df.Age.describe()[5]
      Q3 = df.Age.describe()[6]

      print(Q1,Q2,Q3)

      20.125 28.0 38.0
```

```
[44]: Q1 = df.Age.quantile(0.25)
      Q2 = df.Age.quantile(0.5)
      Q3 = df.Age.quantile(0.75)

      print(Q1,Q2,Q3)

      20.125 28.0 38.0
```

```
[45]: IQR = Q3-Q1
      print(IQR)

      lower_bound = Q1 - 1.5*IQR
      upper_bound = Q3 + 1.5*IQR

      print(lower_bound,upper_bound)

      17.875
      -6.6875 64.8125
```

```
[46]: df.loc[(df.Age>upper_bound) | (df.Age<lower_bound)]['Age']
```

```
[46]: 33      66.0
      54      65.0
      96      71.0
      116     70.5
      280     65.0
```


```
[47]: outlier_indices = df.loc[(df.Age>upper_bound) | (df.Age<lower_bound)].index
      outlier_indices
```

```
[47]: Index([33, 54, 96, 116, 280, 456, 493, 630, 672, 745, 851], dtype='int64')
```

detecting outliers using boxplot


```
[48]: sns.boxplot(df.Age)
```

```
[48]: <Axes: ylabel='Age'>
```



```
[49]: sns.kdeplot(df.Age)
```

```
[49]: <Axes: xlabel='Age', ylabel='Density'>
```



dropping the outliers

```
[50]: df.loc[~((df.Age>upper_bound) | (df.Age<lower_bound))]
```

```
[50]:
```

	PassengerId	Survived	Pclass	Name	Gender	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C

2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S
...
885	886	0	3	Rice, Mrs. William (Margaret Norton)	female	39.0	0	5	382652	29.1250	NaN	Q
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000	NaN	S
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000	B42	S
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C148	C
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.7500	NaN	Q

703 rows × 12 columns

Imputing for Outliers

```
[51]: median = df.Age.median()
      median
```

```
[51]: 28.0
```

```
[52]: mean_age = df.loc[(df.Age<upper_bound)&(df.Age>10)]['Age'].mean()
      print('mean age prior to removing outliers is ',mean_age)
```

```
df.Age.loc[df.Age>upper_bound] = mean_age
print('mean age after imputing for outliers is ',df.Age.mean())
```


```
mean age prior to removing outliers is 31.556338028169016
mean age after imputing for outliers is 29.11034974553202
```

```
[53]: df.Age.loc[df.Age>upper_bound]
```

```
[53]: Series([], Name: Age, dtype: float64)
```


```
[54]: df.boxplot()
```

```
[54]: <Axes: >
```



```
[56]: df.boxplot(["Age", "Fare"])
```

```
[56]: <Axes: >
```



[]:

